

Exploratory Data Analysis

DS 8015

OUTLINE

- 1 Pandas
- 2 Distribution Analysis
- 3 Correlation identification

Pandas

OBJECT CREATION - 1

```
import numpy as np
import pandas as pd
print(pd.__version__) # print pandas version

#Creating a Series by passing a list of values,
# letting pandas create a default integer index:
s = pd.Series([1,3,5,np.nan,6,8])

#Creating a DataFrame by passing a NumPy array,
# with a datetime index and labeled columns:
dates = pd.date_range('20130101', periods=6)
df = pd.DataFrame(np.random.randn(6, 4), index=dates,
                  columns=list('ABCD'))
```

OBJECT CREATION - 2

#Creating a DataFrame by passing a dict of objects

that can be converted to series-like.

```
df2 = pd.DataFrame({  
    'A': np.random.randint(0,10, size=5),  
    'B': pd.Series(2, index=list(range(5)), dtype='float32'),  
    'C': np.array([3]*5, dtype='int32'),  
    'D': pd.Categorical(["test", "train", "test", "train", "train"]),  
    'E': 'foo',  
    'F': pd.date_range(end='20131231', periods=5)})
```

set index of the dataframe

```
df2.index = pd.date_range(start='20130101', periods=5)
```

```
print(df2.dtypes)
```

```
# A int64
```

```
# B float32
```

```
# C int32
```

```
# D category
```

```
# E object
```

```
# F datetime64[ns]
```

VIEWING DATA

```
# View the top/bottom rows of the frame:
```

```
df.head()
```

```
df.tail()
```

```
# Display the index, columns:
```

```
df.index
```

```
df.columns
```

```
# DataFrame.to_numpy() gives a NumPy representation of the data.
```

```
df.to_numpy()
```

```
# Show a quick statistic summary of your (numerical) data:
```

```
df.describe() # df.describe(include="all")
```

```
# transposing the data
```

```
df.T
```

```
# Sorting
```

```
df.sort_index(axis=1, ascending=False) # by axis
```

```
df.sort_values(by='B') # by values
```

SELECTION

```
#We will see optimized pandas data access methods
# .at, .iat, .loc and .iloc.

# Selecting a single column, which yields a Series:
df['A']
# df.A

# Selecting via [], which slices the rows.
df[0:3]
df['20130102':'20130104']
```

SELECTION BY LABEL

```
# For getting a cross section using a label:
```

```
df.loc[dates[0]]
```

```
# Selecting on a multi-axis by label:
```

```
df.loc[:, ['A', 'B']]
```

```
# Showing label slicing, both endpoints are included:
```

```
df.loc['20130102':'20130104', ['A', 'B']]
```

```
# Reduction in the dimensions of the returned object:
```

```
df.loc['20130102', ['A', 'B']]
```

```
# For getting a scalar value:
```

```
df.loc[dates[0], 'A']
```

```
#For getting fast access to a scalar
```

```
# (equivalent to the prior method):
```

```
df.at[dates[0], 'A']
```


SELECTION BY POSITION

Select via the position of the passed integers:

```
df.iloc[3]  
df.iloc[3][1]  
df.iloc[3]['B']
```

By integer slices, acting similar to numpy/python:

```
df.iloc[3:5, 0:2]
```

By lists of int pos. loc., similar to the numpy/python style:

```
df.iloc[[1, 2, 4], [0, 2]]
```

For slicing rows/columns explicitly:

```
df.iloc[1:3, :]  
df.iloc[:, 1:3]
```

For getting a value explicitly:

```
df.iloc[1, 1]
```

For getting fast access to a scalar (equiv. to prior method):

```
df.iat[1, 1]
```

BOOLEAN INDEXING

```
# Using a single column's values to select data.
```

```
df[df.A > 0]
```

```
#Selecting values from a DataFrame
```

```
# where a boolean condition is met.
```

```
df[df > 0]
```

```
# Using the isin() method for filtering:
```

```
df2 = df.copy()
```

```
df2['E'] = ['one', 'one', 'two', 'three', 'four', 'three']
```

```
df2[df2['E'].isin(['two', 'four'])]
```

SETTING

```
# Setting a new column automatically aligns data by indexes.
```

```
s1 = pd.Series([1, 2, 3, 4, 5, 6],  
               index=pd.date_range('20130102', periods=6))  
df['F'] = s1
```

```
# Setting values by label:
```

```
df.at[dates[0], 'A']=0
```

```
# Setting values by position:
```

```
df.iat[0,1]=0
```

```
# Setting by assigning with a NumPy array:
```

```
df.loc[:, 'D'] = np.array([5] * len(df))
```

```
# A where operation with setting.
```

```
df2 = df.copy()  
df2[df2 > 0] = -df2
```

MISSING DATA

- pandas primarily uses the value `np.nan` to represent missing data.
- It is by default not included in computations.
- Reindexing allows you to change/add/delete the index on a specified axis. This returns a copy of the data.

```
df1 = df.reindex(index=dates[0:4], columns=list(df.columns)+['E'])  
df1.loc[dates[0]:dates[1], 'E'] = 1
```

```
# To drop any rows that have missing data.
```

```
df1.dropna(how='any')
```

```
# Filling missing data.
```

```
df1.fillna(value=10)
```

```
# To get the boolean mask where values are nan.
```

```
pd.isna(df1)
```

OPERATIONS - 1

- Operations in general exclude missing data.

```
# Performing a descriptive statistic:
```

```
df.mean()
```

```
# Same operation on the other axis:
```

```
df.mean(1)
```

```
#Subtraction of dataframe and other,  
# element-wise (binary operator sub).
```

```
s = pd.Series([1, 3, 5, np.nan, 6, 8], index=dates).shift(2)  
df.sub(s, axis='index')
```

```
# applying functions to the data
```

```
df.apply(np.cumsum)
```

```
df.apply(np.cumsum, axis=1)
```

```
df.apply(lambda x: x.max() - x.min())
```

OPERATIONS - 2

```
# histogramming
```

```
s = pd.Series(np.random.randint(0, 7, size=10))  
s.value_counts()  
s.hist() # bins=10
```

```
# string methods
```

```
s = pd.Series(['A', 'B', 'C', 'Aaba', 'Baca', np.nan,  
              'CABA', 'dog', 'cat'])  
s.str.lower()
```

```
# Concatenating pandas objects together with concat():
```

```
df = pd.DataFrame(np.random.randn(10, 4))  
# break into pieces  
pieces = [df[:3], df[3:7], df[7:]]  
pd.concat(pieces)
```

OPERATIONS - 3

```
# join: for SQL style merges
```

```
left = pd.DataFrame({'key': ['foo', 'foo'], 'lval': [1, 2]})  
right = pd.DataFrame({'key': ['foo', 'foo'], 'rval': [4, 5]})  
pd.merge(left, right, on='key')
```

```
left = pd.DataFrame({'key': ['foo', 'bar'], 'lval': [1, 2]})  
right = pd.DataFrame({'key': ['foo', 'bar'], 'rval': [4, 5]})  
pd.merge(left, right, on='key')
```

```
# Append rows to a dataframe
```

```
df = pd.DataFrame(np.random.randn(8, 4),  
                  columns=['A', 'B', 'C', 'D'])  
s = df.iloc[3]  
df.append(s, ignore_index=True)
```

GROUPING

- By “group by” we are referring to a process involving one or more of the following steps:
 - (1) Splitting the data into groups based on some criteria
 - (2) Applying a function to each group independently
 - (3) Combining the results into a data structure

```
df = pd.DataFrame({  
    'A': ['foo', 'bar', 'foo', 'bar', 'foo', 'bar', 'foo', 'foo'],  
    'B': ['one', 'one', 'two', 'three', 'two', 'two', 'one', 'three'],  
    'C': np.random.randn(8),  
    'D': np.random.randn(8)  
})
```

```
#Grouping and then applying the sum()  
# function to the resulting groups.  
df.groupby('A').sum()
```

```
#Grouping by multiple columns forms a hierarchical index,  
# and again we can apply the sum function.  
df.groupby(['A', 'B']).sum()
```


RESHAPING

```
# pivot tables
df = pd.DataFrame({
    'A': ['one', 'one', 'two', 'three'] * 3,
    'B': ['A', 'B', 'C'] * 4,
    'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 2,
    'D': np.random.randn(12),
    'E': np.random.randn(12)
})

pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
```

CATEGORICALS

```
# pandas can include categorical data in a DataFrame
df = pd.DataFrame({"id": [1, 2, 3, 4, 5, 6],
                  "raw_grade": ['a', 'b', 'b', 'a', 'a', 'e']})

# Convert the raw grades to a categorical data type.
df["grade"] = df["raw_grade"].astype("category")

#Rename the categories to more meaningful names
# (assigning to Series.cat.categories is inplace!).
df["grade"].cat.categories = ["very good", "good", "very bad"]

#Reorder categories and simultaneously add missing categories
# (methods under Series .cat return a new Series by default).
df["grade"] = df["grade"].cat.set_categories(
    ["very bad", "bad", "medium", "good", "very good"])

# Sorting is per order in the categories, not lexical order.
df.sort_values(by="grade")

# Grouping by a categorical column also shows empty categories.
df.groupby("grade").size()
```

PLOTTING

```
ts = pd.Series(np.random.randn(1000),  
               index=pd.date_range('1/1/2000', periods=1000))  
ts = ts.cumsum()
```

#On a DataFrame, the plot() method is a convenience to plot
all of the columns with labels:

```
df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index,  
                  columns=['A', 'B', 'C', 'D'])  
df = df.cumsum()  
plt.figure()  
plt.legend(loc='best')  
  
df.plot()
```

GETTING DATA IN/OUT

```
# Writing/reading to/from a csv file.
```

```
df.to_csv('foo.csv')
```

```
pd.read_csv('foo.csv')
```

```
# Writing/reading to/from an excel file.
```

```
import openpyxl
```

```
df.to_excel('foo.xlsx', sheet_name='Sheet1')
```

```
pd.read_excel('foo.xlsx', 'Sheet1', index_col=None,  
             na_values=['NA'])
```

Distribution Analysis

VISUAL CHECK - 1

#Reading the dataset in a dataframe using Pandas

```
df = pd.read_csv("LoanPrediction/train.csv")
```

categorical cols

```
df.select_dtypes(exclude=['int', 'float']).columns
```

```
# Index(['Loan_ID', 'Gender', 'Married', 'Dependents',
```

```
# 'Education', 'Self_Employed', 'ApplicantIncome', 'Property_Area',
```

```
# 'Loan_Status'], dtype='object')
```

numerical cols

```
df.select_dtypes(include=['int', 'float']).columns
```

```
# Index(['CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term',
```

```
# 'Credit_History'], dtype='object')
```

#For the non-numerical values, look at frequency distribution

using following command:

```
df['Property_Area'].value_counts()
```

```
# Semiurban 233, Urban 202, Rural 179
```

VISUAL CHECK - 2

```
#can be used for getting various info about data
# (e.g. existence of outliers)

# plotting the histogram of ApplicantIncome
df['ApplicantIncome'].hist(bins=50)

# look at box plots to understand the distributions
df.boxplot(column='ApplicantIncome')
# segregate ApplicantIncome by Education
df.boxplot(column='ApplicantIncome', by = 'Education')
```

CATEGORICAL VARIABLE ANALYSIS

```
# Show the counts of observations in each categorical bin.
import seaborn as sns
sns.countplot(x="Credit_History", hue="Loan_Status",
              data=df, palette="Greys")

# Compute a simple cross-tabulation of two (or more) factors.
pd.crosstab(index=df['Credit_History'],
             columns=df['Loan_Status'], margins=False)

pd.crosstab(index=[df['Credit_History'], df['Education']],
             columns=df['Loan_Status'], margins=False)

# Stacked chart
temp3 = pd.crosstab(df['Credit_History'], df['Loan_Status'])
temp3.plot(kind='bar', stacked=True, color=['red', 'blue'],
           grid=False)
```


TREATING EXTREME VALUES

```
# One approach is to remove rows with outliers
df_red = df
percentile_val = df_red.LoanAmount.quantile(0.975)
# reduce df by removing rows based on percentile_val
df_red = df_red[df_red['LoanAmount'] < percentile_val]

# instead of treating extreme values as outliers, use log
df['LoanAmount_log'] = np.log(df['LoanAmount'])

# scale back the extreme values
df["LoanAmount_new"] = df["LoanAmount"]
df["LoanAmount_new"][df["LoanAmount_new"]>percentile_val]
    = percentile_val

# what happens if you do the following?
df[df["LoanAmount_new"]>percentile_val] = percentile_val
```

Correlation identification

PREPROCESSING AND PLOTTING - 1

```
#many correlation functions such as pearson
# and spearman tests require numerical features
# so, first convert categorical feats to numeric ones

# Married is ['No', 'Yes']
df['Married'] =df['Married'].astype('category').cat.codes
# Now, married is [0, 1]

# check the scatter plots
df.plot.scatter(x = 'ApplicantIncome',
               y = 'CoapplicantIncome', c='DarkBlue')
df.plot.scatter(x = 'ApplicantIncome',
               y = 'CoapplicantIncome', c='DarkBlue', , s=df['Dependents'])
```

PREPROCESSING AND PLOTTING - 2

```
# combined scatter plot matrix
from pandas.plotting import scatter_matrix
attributes = ['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount']
scatter_matrix(df[attributes], figsize=(12, 8))

# seaborn pairplot
df_new = df[['ApplicantIncome', 'CoapplicantIncome',
             'LoanAmount', 'Loan_Status']]
sns.pairplot(df_new, hue='Loan_Status')

# seaborn scatterplot
sns.scatterplot(x="TotalIncome_log", y="LoanAmount_log",
               hue="Loan_Status", data=df, palette="colorblind")
```

CORRELATION TESTS

```
#pearson correlation - used for numeric feats
corr_pearson = df.corr('pearson')
corr_pearson.style.background_gradient(cmap='coolwarm')

#Spearman Rank correlation
# requires ranking the data first
corr_spearman = df.corr('spearman')

#Kendall correlation
# requires ranking the data first
corr_kendall = df.corr('kendall')

#To only obtain the correlation between a
# feature and a subset of the features
df[['ApplicantIncome', 'Education', 'LoanAmount']]
    .corr()['LoanAmount'][:]
```